

Natural Language Processing Assignment

2211CS020372

NEKKANTI KEERTHI

AIML-THETA

1. Correct the Search Query

Explanation: Here is a basic implementation using Python, focusing on spell correction using edit distance and a predefined corpus of words. This code uses `zlib` for compression and `pickle` for serialization, suitable for building an offline model.

Source Code:

1. Correct the Search Query

Explanation: Here is a basic implementation using Python, focusing on spell correction using edit distance and a predefined corpus of words. This code uses `zlib` for compression and `pickle` for serialization, suitable for building an offline model.

Source Code:

```
import re

import pickle

import zlib

from collections import Counter

# Build corpus from a sample dictionary (you can enhance it with more words)
words = """going to china who was the first president of india winner of the match food in america"""

def words_list(text):

    return re.findall(r'\w+', text.lower())

WORDS = Counter(words_list(words))

# Compression for large wordlist

with open('compressed_dict.pkl', 'wb') as f:

    compressed = zlib.compress(pickle.dumps(WORDS))

    f.write(compressed)

# Load dictionary in memory
```

```

def load_dictionary():
    with open('compressed_dict.pkl', 'rb') as f:
        return pickle.loads(zlib.decompress(f.read()))

def edit_distance_one(word):
    letters = 'abcdefghijklmnopqrstuvwxyz'
    splits = [(word[:i], word[i:]) for i in range(len(word) + 1)]
    deletes = [L + R[1:] for L, R in splits if R]
    transposes = [L + R[1] + R[0] + R[2:] for L, R in splits if len(R) > 1]
    replaces = [L + c + R[1:] for L, R in splits if R for c in letters]
    inserts = [L + c + R for L, R in splits for c in letters]
    return set(deletes + transposes + replaces + inserts)

def known(words, dictionary):
    return set(w for w in words if w in dictionary)

def candidates(word, dictionary):
    return (known([word], dictionary) or
            known(edit_distance_one(word), dictionary) or
            [word])

def correct_word(word, dictionary):
    return max(candidates(word, dictionary), key=dictionary.get)

def correct_query(query, dictionary):
    return ''.join(correct_word(word, dictionary) for word in query.split())

# Main correction function

if __name__ == "__main__":
    dictionary = load_dictionary()
    n = int(input())
    queries = [input().strip() for _ in range(n)]
    for query in queries:
        print(correct_query(query, dictionary))

```

2. Deterministic Url and HashTag Segmentation

Explanation: This approach aims to find the most likely and meaningful segmentation of the input strings based on the provided dictionary of words and the constraint of selecting the longest valid tokens from the left.

Source Code:

```
import re

# Load words from words.txt into a set
with open("words.txt", "r") as file:
    dictionary = set(word.strip().lower() for word in file.readlines())

def is_number(s):
    """Check if the string is a number."""
    try:
        float(s)
        return True
    except ValueError:
        return False

def tokenize(input_string, dictionary):
    """
    Tokenize the input string using the longest match first approach.

    Args:
        input_string: The string to be tokenized.
        dictionary: A set of valid words.

    Returns:
```

A list of tokens from the input string.

"""

```
length = len(input_string)
```

```
if length == 0:
```

```
    return []
```

```
# dp[i] stores the tokens for the substring starting from index i
```

```
dp = [None] * (length + 1)
```

```
dp[0] = [] # Base case: empty string has no tokens
```

```
for i in range(1, length + 1):
```

```
    # Consider all possible ending positions for the current substring
```

```
    for j in range(i):
```

```
        left_part = input_string[j:i]
```

```
        # Check if left part is a valid word or number
```

```
        if (left_part in dictionary or is_number(left_part)) and (
```

```
            dp[j] is not None
```

```
        ):
```

```
            # If left part is valid and remaining part has a valid tokenization
```

```
            right_part_tokens = dp[j]
```

```
            right_part_tokens.append(left_part)
```

```
            # Choose the longest valid tokenization
```

```
            if len(right_part_tokens) > len(dp[i]) or dp[i] is None:
```

```
                dp[i] = right_part_tokens
```

```
# Return the tokenization for the entire string if it exists
```

```
return dp[length] if dp[length] is not None else [input_string]
```

```
def main():
```

```

"""Read input strings, tokenize them, and print the results."""

num_test_cases = int(input())

for _ in range(num_test_cases):

    input_string = input().strip().lower()

    # Remove www and extensions for domain names, # for hashtags
    if input_string.startswith("www."):
        input_string = input_string[4:].rsplit(".", 1)[0]
    elif input_string.startswith("#"):
        input_string = input_string[1:]

    tokens = tokenize(input_string, dictionary)

    print(f"Segmentation for Input: {' '.join(tokens)}")


if __name__ == "__main__":
    main()

```

3. Disambiguation: Mouse vs Mouse

Explanation: This code provides a basic framework for classifying the usage of the word "mouse" in a sentence. You can further improve the accuracy by:

- **Expanding the Training Data:** Use a larger and more diverse dataset of sentences.
- **Experimenting with Different Classifiers:** Try other machine learning models like Support Vector Machines (SVM) or Random Forests.
- **Using Word Embeddings:** Consider using word embeddings like Word2Vec or GloVe to capture semantic relationships between words.

Source Code:

```
import pickle

from sklearn.feature_extraction.text import CountVectorizer

from sklearn.naive_bayes import MultinomialNB


# Training data (sample corpus)
training_sentences = [

    "The complete mouse reference genome was sequenced in 2002.",
    "Tail length varies according to the environmental temperature of the mouse "
    "during postnatal development.",
    "A mouse is an input device.",
    "Many mice have a pink tail.",
    "The mouse pointer on the screen helps in navigation.",
    "A rodent like a mouse has sharp teeth.",
    "The mouse was connected to the computer using a USB port.",
    "The house was infested with mice.",
    "Computer users often prefer a wireless mouse."

]


# Labels corresponding to the training sentences
```

```

labels = [
    "animal",
    "animal",
    "computer-mouse",
    "animal",
    "computer-mouse",
    "animal",
    "computer-mouse",
    "animal",
    "computer-mouse"
]

```

```

# Vectorize the training sentences

```

```

vectorizer = CountVectorizer()

```

```

X_train = vectorizer.fit_transform(training_sentences)

```

```

# Create and train the Naive Bayes classifier

```

```

classifier = MultinomialNB()

```

```

classifier.fit(X_train, labels)

```

```

# Function to predict the type of mouse

```

```

def predict_mouse_type(sentence):

```

```

    """

```

```

    Predicts whether the 'mouse' in the sentence refers to an animal or a computer mouse.

```

```

    Args:

```

```

        sentence: The input sentence.

```

```

    Returns:

```

```

        "animal" or "computer-mouse"

```

```

    """

```

```
vectorized_sentence = vectorizer.transform([sentence])  
prediction = classifier.predict(vectorized_sentence)[0]  
return prediction
```

```
# Get number of test cases
```

```
num_test_cases = int(input())
```

```
# Process each test case
```

```
for _ in range(num_test_cases):
```

```
    sentence = input()
```

```
    prediction = predict_mouse_type(sentence)
```

```
    print(prediction)
```

```
# Optionally, save the trained model for later use
```

```
with open('mouse_classifier.pkl', 'wb') as f:
```

```
    pickle.dump((vectorizer, classifier), f)
```


4. Language Detection

- **Explanation:** This function loads the pre-trained model from a serialized file.
- It takes a text snippet as input, normalizes it to ASCII, and converts it into a TF-IDF vector using the loaded vectorizer.
- The function then uses the trained classifier to predict the language of the snippet based on the extracted features.

Source Code:

```
import pickle

import unicodedata

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB


def normalize_to_ascii(text):

    """Remove non-ASCII characters and normalize text."""

    return unicodedata.normalize("NFKD", text).encode("ascii", "ignore").decode("ascii")


# Step 1: Training Data

training_texts = {

    "English": [

        "The quick brown fox jumps over the lazy dog.",

        "Rip Van Winkle is a story set in the years before the American Revolutionary War.",

    ],

    "French": [

        "Le renard brun rapide saute par-dessus le chien paresseux.",

        "La revolution francaise a marque une periode importante de l'histoire.",

    ],

    "German": [

        "Der schnelle braune Fuchs springt uber den faulen Hund.",

        "Die deutsche Wiedervereinigung war ein historisches Ereignis.",

    ],
```

```

"Spanish": [
    "El rapido zorro marron salta sobre el perro perezoso.",
    "La Revolucion Espanola fue un momento clave en la historia."
    "Si quieres que te asciendan te tienes que poner las pilas.",
],
}

```

```

# Normalize training data to ASCII

```

```

labels = []

```

```

texts = []

```

```

for language, samples in training_texts.items():

```

```

    labels.extend([language] * len(samples))

```

```

    texts.extend([normalize_to_ascii(sample) for sample in samples])

```

```

# Step 2: Preprocessing and Feature Extraction

```

```

vectorizer = TfidfVectorizer(ngram_range=(2, 4), analyzer="char")

```

```

X_train = vectorizer.fit_transform(texts)

```

```

# Step 3: Train the Model

```

```

classifier = MultinomialNB()

```

```

classifier.fit(X_train, labels)

```

```

# Step 4: Serialize the Model

```

```

with open("language_model.pkl", "wb") as model_file:

```

```

    pickle.dump((vectorizer, classifier), model_file)

```

```

# Step 5: Language Detection Function

```

```

def detect_language(snippet):

```

```

    with open("language_model.pkl", "rb") as model_file:

```

```

        vectorizer, classifier = pickle.load(model_file)

```

```
# Normalize snippet to ASCII

snippet = normalize_to_ascii(snippet)

X_test = vectorizer.transform([snippet])

prediction = classifier.predict(X_test)

return prediction[0]
```

```
# Input Processing
```

```
if __name__ == "__main__":
```

```
    # Read multi-line input
```

```
    snippet = ""
```

```
    while True:
```

```
        try:
```

```
            line = input()
```

```
            if line.strip():
```

```
                snippet += line + "\n"
```

```
        except EOFError:
```

```
            break
```

```
# Predict and Output
```

```
detected_language = detect_language(snippet.strip())
```

```
print(detected_language)
```

5. The Missing Apostrophes

Explanation Apostrophe Handling: The code defines a function `restore_apostrophes` that iterates through each word in the input text. It uses a combination of explicit checks for common contractions (e.g., "don't," "can't," "I've") and a regular expression to handle possessive nouns (e.g., "cat's," "dog's") to restore apostrophes where appropriate.

Source Code:

```
import re

# Function to handle apostrophes for contractions and possessives
def restore_apostrophes( text):
    restored_text = []
    words = text.split()

    for word in words:
        lower_word = word.lower()

        # Handle contractions
        if word.lower() == "dont":
            restored_text.append("don't")
        elif word.lower() == "wont":
            restored_text.append("won't")
        elif word.lower() == "cant":
            restored_text.append("can't")
        elif word.lower() == "isnt":
            restored_text.append("isn't")
        elif word.lower() == "arent":
            restored_text.append("aren't")
        elif word.lower() == "wasnt":
            restored_text.append("wasn't")
```

```
elif word.lower() == "werent":
    restored_text.append("weren't")
elif word.lower() == "hasnt":
    restored_text.append("hasn't")
elif word.lower() == "havent":
    restored_text.append("haven't")
elif word.lower() == "hadnt":
    restored_text.append("hadn't")
elif word.lower() == "didnt":
    restored_text.append("didn't")
elif word.lower() == "ive":
    restored_text.append("I've")
elif word.lower() == "were":
    restored_text.append("we're")
elif word.lower() == "i":
    restored_text.append("I")
elif word.lower() == "id":
    restored_text.append("I'd")
elif word.lower() == "ive":
    restored_text.append("I've")
elif word.lower() == "youve":
    restored_text.append("you've")
elif word.lower() == "hes":
    restored_text.append("he's")
elif word.lower() == "shes":
    restored_text.append("she's")
elif word.lower() == "its":
    restored_text.append("it's")
elif word.lower() == "were":
    restored_text.append("we're")
```

```
# Handle possessives (only add 's when it makes sense)
```

```
elif re.match(r'\w+s$', word) and lower_word not in ["its", "hers", "ours", "yours", "theirs"]:
```

```
    restored_text.append(re.sub(r"s$", "'s", word))
```

```
# For normal words that don't need apostrophes, keep them as is
```

```
else:
```

```
    restored_text.append(word)
```

```
return " ".join(restored_text)
```

```
# Input
```

```
input_text = ""At a news conference Thursday at the Russian manned-space facility in Baikonur, Kazakhstan, Kornienko said "we will be missing nature, we will be missing landscapes, woods." He admitted that on his previous trip into space in 2010 "I even asked our psychological support folks to send me a calendar with photographs of nature, of rivers, of woods, of lakes."
```

```
Kelly was asked if hed miss his twin brother Mark, who also was an astronaut.
```

```
"Were used to this kind of thing," he said. "Ive gone longer without seeing him and it was great."
```

```
The mission wont be the longest time that a human has spent in space - four Russians spent a year or more aboard the Soviet-built Mir space station in the 1990s.
```

```
SCI Astronaut Twins
```

```
Scott Kelly (left) was asked Thursday if hed miss his twin brother, Mark, who also was an astronaut. Were used to this kind of thing, he said. Ive gone longer without seeing him and it was great. (NASA/Associated Press)
```

```
"The last time we had such a long duration flight was almost 20 years and of course al{-truncated-}"
```

```
# Restore apostrophes
```

```
output_text = restore_apostrophes(input_text)
```

```
print(output_text)
```

6. Segment the Twitter Hashtags

Explanation: Tokenization with Dynamic Programming: The `segment_hashtag` function uses dynamic programming to break down the hashtag into a sequence of words. It iterates through the hashtag, checking for valid word combinations from a given dictionary and selecting the longest possible valid sequence.

Source Code:

```
# Define a function that segments a single hashtag into words
def segment_hashtag(hashtag, word_dict):

    n = len(hashtag)

    dp = [None] * (n + 1)

    dp[0] = [] # Base case: empty string can be segmented as an empty list

    # Iterate over the hashtag string
    for i in range(1, n + 1):

        for j in range(max(0, i - 20), i): # Limit the length of words checked

            word = hashtag[j:i]

            if word in word_dict and dp[j] is not None:

                dp[i] = dp[j] + [word]

                break

    return "".join(dp[n]) if dp[n] is not None else hashtag

# Main function to process input and output results
def process_hashtags(num_hashtags, hashtags, word_dict):

    result = []

    for hashtag in hashtags:

        segmented = segment_hashtag(hashtag, word_dict)

        result.append(segmented)

    return result
```

```
# Sample dictionary of common words (expand this as needed)

word_dict = {
    "we", "are", "the", "people", "mention", "your", "faves",
    "now", "playing", "walking", "dead", "follow", "me"
}


# Sample input

num_hashtags = int(input())

hashtags = [input().strip() for _ in range(num_hashtags)]


# Process the hashtags and print the result

segmented_hashtags = process_hashtags(num_hashtags, hashtags, word_dict)

for segmented in segmented_hashtags:
    print(segmented)
```


7. Expand the Acronyms

Explanation: Acronym Extraction: The code extracts acronyms and their potential expansions from a given set of text snippets by identifying uppercase words within parentheses and searching for preceding phrases. It also attempts to extract acronyms not explicitly defined in parentheses by analyzing the surrounding context.

Source Code:

```
import re

def extract_acronyms_and_expansions(snippets):
    """
    Extract acronyms and their expansions from the provided snippets.
    """
    acronym_dict = {}
    for snippet in snippets:
        # Find all potential acronyms (uppercase words typically enclosed in parentheses)
        matches = re.findall(r'\((\b[A-Z]+\b)\)', snippet)

        for match in matches:
            # Extract the preceding text (potential expansion)
            preceding_text = snippet.split(f'({match})')[0].strip()

            # Look for the last meaningful phrase before the acronym
            expansion_candidates = re.split(r'[.,;:-]', preceding_text)
            if expansion_candidates:
                expansion = expansion_candidates[-1].strip()
                acronym_dict[match] = expansion

        # Additionally, handle acronyms not in parentheses but defined explicitly
        words = snippet.split()
        for i, word in enumerate(words):
```

```

        if word.isupper() and len(word) > 1: # Likely an acronym

            if word not in acronym_dict:

                # Try to extract its expansion from the surrounding context

                if i > 0:

                    preceding_context = " ".join(words[max(0, i-5):i])

                    if preceding_context:

                        acronym_dict[word] = preceding_context

            return acronym_dict

def process_tests(acronym_dict, tests):
    """
    Process test acronyms and return their expansions.
    """
    results = []
    for test in tests:
        # Normalize the test acronym (case insensitive)
        expansion = acronym_dict.get(test.upper(), "Not Found")
        results.append(expansion)
    return results

def main():
    # Read input
    n = int(input().strip())
    snippets = [input().strip() for _ in range(n)]
    tests = [input().strip() for _ in range(n)]

    # Extract acronyms and expansions
    acronym_dict = extract_acronyms_and_expansions(snippets)

    # Process test queries
    results = process_tests(acronym_dict, tests)

    # Output results

```

```
print("\n".join(results))  
if __name__ == "__main__":main()
```

8. Correct the Search Query

Explanation: Here is a basic implementation using Python, focusing on spell correction using edit distance and a predefined corpus of words. This code uses `zlib` for compression and `pickle` for serialization, suitable for building an offline model.

Source Code:

```
import re

import pickle

import zlib

from collections import Counter


# Build corpus from a sample dictionary (you can enhance it with more words)

words = """going to china who was the first president of india winner of the match food in america"""


def words_list(text):

    return re.findall(r'\w+', text.lower())


WORDS = Counter(words_list(words))


# Compression for large wordlist

with open('compressed_dict. pkl', 'wb') as f:

    compressed = zlib.compress(pickle.dumps(WORDS))

    f.write(compressed)


# Load dictionary in memory

def load_dictionary():

    with open('compressed_dict. pkl', 'rb') as f:

        return pickle.loads(zlib.decompress(f.read()))


def edit_distance_one(word):

    letters = 'abcdefghijklmnopqrstuvwxyz'

    splits = [(word[:i], word[i:]) for i in range(len(word) + 1)]
```

```

deletes = [L + R[1:] for L, R in splits if R]

transposes = [L + R[1] + R[0] + R[2:] for L, R in splits if len(R) > 1]

replaces = [L + c + R[1:] for L, R in splits if R for c in letters]

inserts = [L + c + R for L, R in splits for c in letters]

return set(deletes + transposes + replaces + inserts)


def known(words, dictionary):

    return set(w for w in words if w in dictionary)


def candidates(word, dictionary):

    return (known([word], dictionary) or

            known(edit_distance_one(word), dictionary) or

            [word])


def correct_word(word, dictionary):

    return max(candidates(word, dictionary), key=dictionary.get)


def correct_query(query, dictionary):

    return ''.join(correct_word(word, dictionary) for word in query.split())


# Main correction function
if __name__ == "__main__":

    dictionary = load_dictionary()

    n = int(input())

    queries = [input().strip() for _ in range(n)]

    for query in queries:

        print(correct_query(query, dictionary))

```

G.A Text-Processing Warmup

Explanation: Article and Date Counting: The code defines a function `count_articles_and_dates` that takes a text fragment as input. It first normalizes the text to lowercase for case-insensitive article counting. Then, it uses regular expressions to count occurrences of the definite and indefinite articles ("a," "an," "the") and identify valid dates in various formats (e.g., "DD Month YYYY," "Month DD, YYYY," etc.)

Source Code:

```
import re

def count_articles_and_dates(fragment):
    """
    Count occurrences of 'a', 'an', 'the', and valid dates in a given text fragment.
    """
    # Normalize text for article counting
    lower_fragment = fragment.lower()

    # Count articles
    a_count = len(re.findall(r'\b[a]\b', lower_fragment))
    an_count = len(re.findall(r'\b[an]\b', lower_fragment))
    the_count = len(re.findall(r'\b[the]\b', lower_fragment))

    # Identify valid dates
    date_patterns = [
        r'\b\d{1,2}?(?:st|nd|rd|th)?(?:\s+of)?\s+(January|February|March|April|May|June|July|August|September|October|November|December)\s+\d{2,4}\b', # Day Month Year
        r'\b(January|February|March|April|May|June|July|August|September|October|November|December)\s+\d{1,2}?(?:st|nd|rd)?(?:\s+\d{2,4}\b', # Month Day Year
        r'\b\d{1,2}/\d{1,2}/\d{2,4}\b', # Day/Month/Year
        r'\b\d{4}-\d{2}-\d{2}\b' # ISO format: Year-Month-Day
    ]
```

```

# Combine all date patterns

date_regex = '|'.join(date_patterns)

dates = re.findall(date_regex, fragment, re.IGNORECASE)

date_count = len(dates)


return a_count, an_count, the_count, date_count


def main():

    import sys

    input = sys.stdin.read

    # Read input data

    data = input().strip().split("\n")

    t = int(data[0]) # Number of test cases

    fragments = data[1:] # Remaining lines contain the fragments


    results = []

    for i in range(t):

        fragment = fragments[i].strip()

        # Count articles and dates

        a_count, an_count, the_count, date_count = count_articles_and_dates(fragment)

        results.append(f"{a_count}\n{an_count}\n{the_count}\n{ date_count}")


    # Output results

    print("\n".join(results))


if __name__ == "__main__":

    main()

```

10. Who is it?

Explanation: Pronoun Identification and Entity Matching: The code first finds all pronouns (words enclosed in double backslashes) and their positions in the text. It then cleans the text by removing the backslashes. Next, it iterates through each pronoun and searches for the closest matching entity (from a provided list) that appears before the pronoun in the text.

Source Code:

```
import re

def resolve_pronouns(text, entities):

    # Extract all pronouns and their positions

    pronoun_pattern = r'\\(\\w+)\\'

    pronouns = [(match.group(1), match.start()) for match in re.finditer(pronoun_pattern, text)]

    # Clean the text by removing ** markers

    clean_text = re.sub(r'\\(\\w+)\\', r'1', text)

    # Initialize a list to store the resolved entities

    resolved = []

    # For each pronoun, find the corresponding entity

    for pronoun, pos in pronouns:

        closest_entity = None

        closest_distance = float('inf')

        # Iterate through all entities to find the best match

        for entity in entities:

            entity_pos = clean_text.rfind(entity, 0, pos) # Find the last occurrence of the entity before the
            pronoun

            if entity_pos != -1:

                distance = pos - (entity_pos + len(entity))

                if distance < closest_distance:
```



```

        closest_distance = distance

        closest_entity = entity

    # Append the resolved entity to the list
    resolved.append(closest_entity)

return resolved

def main():
    import sys

    input = sys.stdin.read
    data = input().strip().split("\n")

    # Read the number of lines in the text snippet
    n = int(data[0])

    # Combine the next N lines into the full text snippet
    text_snippet = " ".join(data[1:n + 1])

    # Read the list of entities
    entities = [e.strip() for e in data[n + 1].split(';')]

    # Resolve pronouns
    result = resolve_pronouns(text_snippet, entities)

    # Output the resolved entities
    for entity in result:
        print(entity)

```

```
if __name__ == "__main__":  
    main()
```